
Algorithm 1 The ADAM algorithm computes a batch stochastic gradient to compute momentum and RMSProp vectors at each timestep, with a bias correction step accounting for first and second moment estimates. Model parameters are updated using this recursive gradient information.

Require: parameters θ , stochastic objective function $f_i(\theta)$

Require: learning rate η , exponential decay rates $\beta_1, \beta_2 \in [0, 1)$, tolerance ϵ , batch size

Initialize: initial parameter vector θ_0 , initial 1st moment vector $m_0 = 0$, initial 2nd moment vector $v_0 = 0$, initial timestep $t = 0$

while θ_t is not converged **do**

$t = t + 1$

$g_t = \nabla_{\theta} f_t(\theta_{t-1})$ *(batch gradient at iteration t)*

$m_t = \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$ *([Monentum] udpate baised first moment estimate)*

$v_t = \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$ *([RMSProp] udpate baised second raw moment estimate)*

$\hat{m}_t = m_t / (1 - \beta_1^t)$ *([Monentum] bias-corrected first moment estimate)*

$\hat{v}_t = v_t / (1 - \beta_2^t)$ *([RMSProp] bias-corrected second raw moment estimate)*

$\theta_t = \theta_{t-1} - \eta \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$ *([Momentum + RMSProp] update parameters)*

end while

return θ_t

Algorithm 2 SARAH +

Require: objective function $f(\theta)$

Require: learning rate $\alpha > 0$, inner loop size m , outer loop size T

Initialize: initialize arbitrary parameter vector $\tilde{\theta}_0 \in \mathbb{R}$

Iterate:

for $s = 1, 2, \dots, T$ **do**

$\theta_0 = \tilde{\theta}_{s-1}$

$v_0 = \frac{1}{n} \sum_{i=1}^n \nabla f_i(\theta_0)$ *(outer loop full gradient computation)*

$\theta_1 = \theta_0 - \alpha v_0$ *(compute one outer loop parameter update for inner loop)*

$t = 1$

while $\|v_{t-1}\|^2 > \gamma \|v_0\|^2$ **and** $t < m$ **do**

 Sample $i \in \{1, \dots, n\}$ uniformly at random

$v_t = \nabla f_{i_t}(\theta_t) - \nabla f_{i_t}(\theta_{t-1}) + v_{t-1}$ *(gradient estimate (SARAH update))*

$\theta_{t+1} = \theta_t - \alpha v_t$ *(inner loop parameter update)*

$t = t+1$

end for

 Set $\tilde{\theta}_s = \theta_t$

end for

Algorithm 3 SVRG

Require: objective function $f(\theta)$

Require: learning rate $\alpha > 0$, inner loop size m , outer loop size T

Initialize: initialize arbitrary parameter vector $\tilde{\theta}_0 \in \mathbb{R}$

Iterate:

for $s = 1, 2, \dots, T$ **do**

$\tilde{\theta} = \tilde{\theta}_{s-1}$

$\tilde{g} = \frac{1}{n} \sum_{i=1}^n \nabla f_i(\tilde{\theta})$ (*outer loop full gradient computation*)

$\theta_0 = \tilde{\theta}$

for $t = 1, \dots, m$ **do**

 Sample $i \in \{1, \dots, n\}$ uniformly at random

$v_{t-1} = \nabla f_{i_t}(\theta_{t-1}) - \nabla f_{i_t}(\tilde{\theta}) + \tilde{g}$ (*inner loop gradient approximation*)

$\theta_t = \theta_{t-1} - \alpha \cdot v_{t-1}$ (*inner loop parameter update*)

end for

Option 1: Set $\tilde{\theta}_s = \theta_t$ with $t \in \{0, 1, \dots, m\}$ chosen uniformly at random

Option 2: Set $\tilde{\theta}_s = \theta_m$

end for

Algorithm 4 The SARAH algorithm is identical to SVRG except for the *SARAH update*, which modifies the stochastic gradient estimate to use recursive gradient estimate information rather than the initialized gradient to update the gradient estimate in the inner loop.

Require: objective function $f(\theta)$

Require: learning rate $\alpha > 0$, inner loop size m , outer loop size T

Initialize: initialize arbitrary parameter vector $\tilde{\theta}_0 \in \mathbb{R}$

Iterate:

for $s = 1, 2, \dots, T$ **do**

$\theta_0 = \tilde{\theta}_{s-1}$

$v_0 = \frac{1}{n} \sum_{i=1}^n \nabla f_i(\theta_0)$ (*outer loop full gradient computation*)

$\theta_1 = \theta_0 - \alpha v_0$ (*compute one outer loop parameter update for inner loop*)

for $t = 1, \dots, m - 1$ **do**

 Sample $i \in \{1, \dots, n\}$ uniformly at random

$v_t = \nabla f_{i_t}(\theta_t) - \nabla f_{i_t}(\theta_{t-1}) + v_{t-1}$ (*gradient estimate (SARAH update)*)

$\theta_{t+1} = \theta_t - \alpha v_t$ (*inner loop parameter update*)

end for

 Set $\tilde{\theta}_s = \theta_t$ with t chosen uniformly at random from $\{0, 1, \dots, m\}$

end for
